


Fair CV screening, from raw data to a live app



THE MODEL

all-MiniLM-L6-v2

- 6-layer sentence encoder, 384-d vectors
- CVs cut into 180-word chunks, averaged
- PASS if cosine ≥ 0.436 (AUC 0.78)

THE DEBIAS LAYER

LEACE concept erasure

- Closed form: $x' = x \cdot A + b$
- Fitted on 36 cities $\times$ 2 genders
- Belrose et al., 2023

LOCATION DECISION FLIPS

15.3% → 10.4%

Before → after debias. Matching AUC kept:
0.784 → 0.780

Built once offline, served by one ONNX model

